

1. Abstract

hej med dig

- The tool can be found here: <https://github.com/dag-andersen/kubernetes-reference-highlighter>
- and downloaded here: <https://github.com/dag-andersen/kubernetes-reference-highlighter/releases>
- or downloaded from the VS Code Marketplace.

2. CCS CONCEPTS

- Human-centered computing → Open source software; Software Engineering Tooling → Code Validation

3. Keywords

- Open Source Software, IDE Extension, Code Linting, Static code analysis, Code Validation, Kubernetes

4. Introduction

Container Orchestration and Kubernetes are continuously getting more popular during the last couple of years.

Around 25-30% of professional developers use Kubernetes (25.45% [[StackOverflow Survey 2022](#)], 30% [[SlashData's Developer Economics survey](#)], and around one fourth of the remaining wats to work with Kubernetes in the future [[StackOverflow Survey 2022](#)]. So we can expect that more and more developer will be deploying and configuring infrastructure on kubernetes through YAML.

But Kubernetes is known for being remarkably complex and the learning curve can be quite steep. In a survey done by StackOverflow 74,75% of the developers who use Kubernetes "loves it" and the remaining 25,25 "dreads it" [\[source\]](#). Why the ~25% dreads Kubernetes is not specified in the survey, but anything that would lower that percentage is worth striving for.

The eco-systems around Kubernetes is in constant development, and it can be difficult to keep up with the newest tools and features. When building workload aimed at running on Kubernetes most configuration is down in YAML and revolve around Kubernetes Object Definitions. This means we as community need to ensure we have proper tooling for validating this infrastructure. Both to speed up the development process but also to reduce the amount of errors that goes into production. The sooner the bug is found in the software process the cheaper it is to correct [\[source\]](#).

5. Problem

It is very easy to make mistakes when when deploying to Kubernetes. There is strict rules of what fields you can add in your Kubernetes Object Definition, but there is no checks that the values provided are valid (besides the fact it needs to be the correct type) - e.g. your *Pod* can request environments variables from a secret that does not exist, or your *service* can point to Pods that does not exist.

Terraform does static "compile" time checks, so you will be notified about your broken references beforehand. By design Kubernetes does not do that and can not do that. Kubernetes objects are not expected to be created at the same time or in a given order, so we can not talk about a "compile" time. Kubernetes is managed with Control Planes and there is no order to when which resource is created/scheduled. In kubernetes your references will often not exist on creation time, but will only be created at a later point. This explains why there is no static validation by default when creating Kubernetes resources.

Kubernetes objects definition are written in YAML. YAML is a data serialization language, that only knows the concepts of primitive data types like arrays, strings, numbers, etc and therefore does not know the concept of references. All references declared in your Kubernetes objects definition is written as a string, with no type checking or validation. This means there is no built-in reference validation in YAML.

Since most IDEs, Plugin, or Tools do not verify any of the magic strings, it can quite cumbersome to debug the code. It tend to involve manually reading the magic string repeatably until you discover that there is typo, or each of your resources exists in two different namespaces and therefor cant communicate. No public and free IDE feature or extension exists that checks this.

Furthermore, developers often use tools like *Helm* and *Kustomize* to template their YAML-files. This is done so multiple configuration can inherit/share common code across different configuration. This means that references that is hardcoded in the files may not exist in any of the plain files, but may only be generated on `runtime` when one of the templating tools are used. This makes it much harder to give valuable information to the developer because kubernetes object's names are dynamically generated by the tool. I have not found any extension that tries to tackle this challenge, so currently developer doesn't get any assistance validating their YAML-files live while coding when using templating tools like *Helm* and *Kustomize*.

In dynamic languages like JavaScript, where there is no *compile time*-checks either, it has been shown that developers use static analysis tools such (or so called *Linters*) like *ESLint* to prevent errors. Preventing Errors is said to be the number one reason to use a linter, because it catches the bugs early on, so you don't have to spend time debugging it on runtime [\[source\]](#). Linters like *ESLint* can be integrated in popular IDEs like VS Code or IntelliJ IDEA [\[source\]](#). The linter gives the developer continuous feedback on if their code have errors or not. So even though a language does not have static compile time checks, it is still possible to reduce a substantial amount of errors by using Linters running in your IDE.

Currently plenty of open source validation tools exists for Kubernetes Object Definition, but most of them are Command Line tools and does not give continuously feedback live while the developer is coding. Furthermore the tools only validate the definitions based on the local files provided to the tool and not what is actually running in a Kubernetes Cluster. Most often than not you only have a subset of the full infrastructure configuration on your local machine, so many of the code references and objects will naturally be broken and the tools will give incorrect/adequate valuation. So in order to overcome this we need to look outside the borders of the IDE/current folder. (This ofc only works if you have read access to those Kubernetes endpoints/resources from you local machine. If you the endpoints is only visible from inside an enclosed environment this will not work.)

This paper will showcase a prototype of an extension made for visual studio that validates object-references in Kubernetes Object Definitions, continuously while the developer is typing. The tool aims to incorporate both plain local YAML-files, output generated by tools like Helm or Kustomize, and scan what resources already exist in a running Kubernetes Cluster. The aim is to provide the developer with valuable information that reduce the number of bugs encountered when deploying and speed up the debugging process. I will refer to the extension as *Kubernetes Reference Highlighter* (KRH).

Ten developers are given a task to debug a system. Half of them will do it with the extension enabled and the others won't. The time it takes for each developer will be tracked and evaluated upon.

6. Related tools | Existing tools

In this section existing tools are listed that aim to help the developer create correct/valid Kubernetes Object Definitions and tackle some of the challenges explained in the introduction-section.

- **JetBrains IntelliJ IDEA Kubernetes Plugin** [[link](#)]: The full version IntelliJ that includes support for plugins is a paid product and therefore is not accessible publicly. JetBrains' only takes local files into account when highlighting references. The plugin does not check your current cluster or build anything with kustomize like KRH. But JetBrains' plugin comes with extra functionality that lets the user jump between files by clicking the highlighted references (something that KRH does not).
- **KubeLinter** [[docs](#)] [[GitHub](#)]: This is an open source Command Line Tool that validates your Kubernetes manifests. KubeLinter is a feature-rich tool that informs and warns of all sorts of issues that may be in their code. The full list of features can be found [[Here](#)]. KubeLinter is ideal to integrate in testing-pipelines where it can function as an initial step for everything YAML before it goes into production. KubeLinter can validate Helm Charts, but can not validate Kustomize templates. The tool only checks the local files it is provided, so it does not know what already exists in a cluster.
- **kube-score** [[GitHub](#)] [[Validation Checks](#)]: kube-score is a tool that performs static code analysis of your Kubernetes object definitions. Even though there are over 30 different checks the values are quite limited. The tool seems to only validate all the Kubernetes object definitions individually, and not as a whole (with e.g. reference checking). Furthermore, just like *KubeLinter* it is a Command Line tool and therefore can be integrated as part of a testing-pipeline, and does not provide any live feedback when coding in an IDE.
- **Kubevious** [[docs](#)] [[GitHub](#)]. This tool comes with 44 built-in Kubernetes Resource validation checks. The checks include cross-file checks that e.g. validate if a `service` points to two deployments at the same time and that sort of misconfiguration. Kubevious has a powerful toolset just like KubeLinter, but this tool has a UI and needs to be run in a cluster or as a standalone program. It can only run its checks on resources already deployed to Kubernetes, and therefore does not do any live validation for the developer while coding. This tool is mainly for finding and debugging issues after it has already been deployed.
- **Conftest** [[GitHub](#)], **Copper** [[GitHub](#)], and **Config-lint** [[GitHub](#)] are all similar tools that try to validate Kubernetes Manifests/YAML, but none of them come with built-in checks, so the user has to make their own "rules" that the YAML will be validated against. All these tools are Command Line tools, so they do not provide any live feedback but only validation when command run. These tools accelerate when you want to build your own custom checks (maybe on your own Custom Resource Definition (CRDs)).

What is shared across all the tools listed above is that none of them manage to do references checking across multiple files in a continuous manner that both takes local files, templating tools, and actual running clusters into account, and being free to use at the same time.

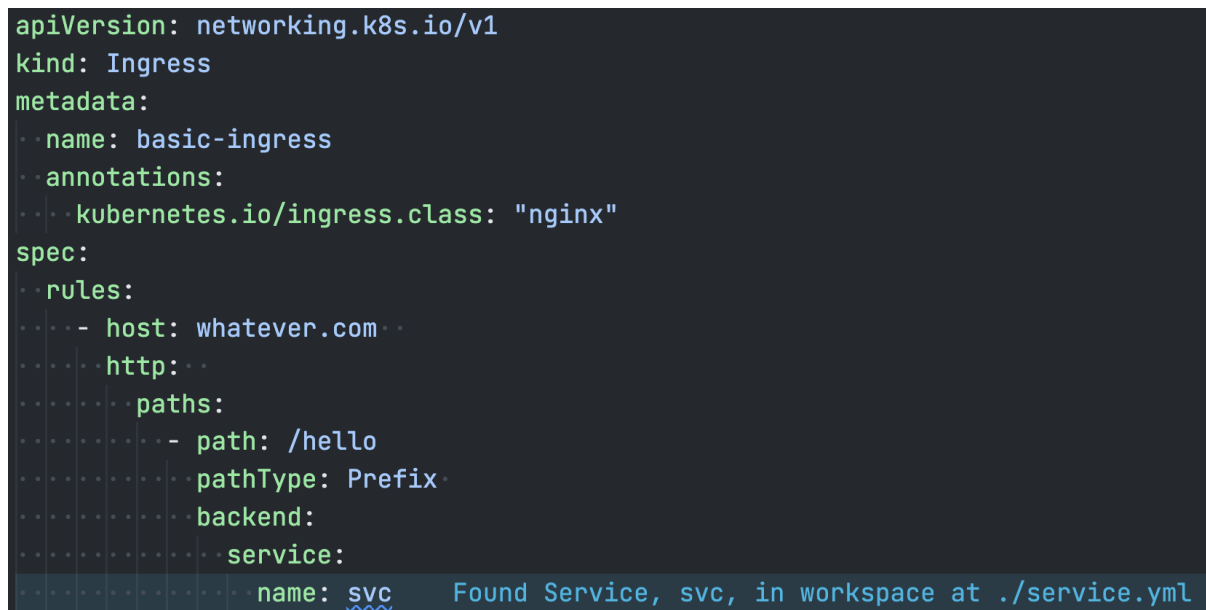
7. Prototype | Implementation

This section will describe a prototype of the *Kubernetes Reference Highlighter* VS Code Extension.

7.1. Overview

The tool is built as a Visual Studio Code (VS Code) extension, since it is free and is the most popular IDE/Editor in the industry [source]. The extension is made in typescript, since that is the default for developing VS Code Extensions, since VS Code is written in typescript.

The extension highlights the name of an object if it finds a reference in the open file. The highlighting updates in real time while you type, giving you constant feedback on if your references are found or not.



```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: basic-ingress
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
    - host: whatever.com
      http:
        paths:
          - path: /hello
            pathType: Prefix
            backend:
              service:
                name: svc Found Service, svc, in workspace at ./service.yml
```

Fig. 1 shows how the extension would highlight the string `svc` in the VS Code because it found an object with that name file `./service.yml`. On this image the VS Code extension, *Error Lens*, is enabled to better visualize the highlighting.

A tool like this is close to useless if the developer can not rely on validity of the feedback it gives. This is why this tool aims to only highlight a string if it is sure that the reference exists. It is better to have false negatives than false positives in this case.

This extension make use of `vscode.diagnostics` for highlighting strings. Diagnostics are the squiggly line under code lines, and the severity can either be marked as an `Severity.Error` , `Severity.Warning` , `Severity.Information` , or `Severity.Hint` . This extension use the `Severity.Information` for all object-references, and it only use `Severity.Error` , when `Kustomize build <path>` fails. Since these Diagnostics are built into the VS Code IDE, then other extensions or tools can also read and act upon the information provided by this extension. This makes this extension integrate well with other VS Code features/extensions. In particular i suggest installing *Error Lens* together with this extension to visualize the diagnostics even better for the user.

The extensions architectural structure can be split up into two components/steps: *Reference Collection* and *Reference Detection*. *Reference Collection* is the step that collects and builds a list of all the objects the extension can find. This step is triggered every time you save a YAML-file in the workspace. *Reference Detection* is the step that parses the current open file and try to find references that matches objects found in the *Reference Collection*-step. This step is triggered every time a YAML-file is changed in the workspace.

7.2. Reference Detection

As of Version 0.0.2 the extension highlights references to `Services` , `Deployments` , `Secrets` , and `ConfigMaps` . More kinds can easily be added in the future.

The extension use YAML-parsing and Regex as the main method for reading the files in the workspace.

All resources are namespace-sensitive. An object-reference will not be highlighted if the resource exclusively exists in another namespace.

Raw hostname addresses pointing to `Service` s are a bit different and the most difficult one to handle. You can either call a `Service` by using its name as an address if you call it from inside its namespace (`my-service`) - or you can call it from anywhere within the cluster by suffixing the `Service` -name with its namespace (`my-service.its-namespace`). This tool will only highlight an address pointing to a service, if it exist inside the same namespace or the address is suffixed with its namespace. This is an important distinction that can prevent many issues related to traffic routing, where you try to access a service which may be in a different namespace.

Raw address are en general difficult to detect, because it can be written as part of a string in many places. E.g. the address of an service can be written as part of an environment variable accessible from a Pod or it can be stored as a field in a `Configmaps`, which is read from a Pod. The change of detecting these references are more challenges, and there is a bigger risk of detecting false positives.

If the string matches with multiple objects the extension will show all the references. A reference can as an example both exists in the running cluster and in local file, because the object is already deployed to the cluster.

< insert picture >

KubeLinter detects "dangling references", so it inform the user that a references does not exist. This is great, but often times you do not have the whole infrastructure configuration locally, meaning that the dangling references that it detects may not be dangling after all, because the references actually exists in the cluster. This ensures the KubeLinter can report false positives (detecting an issue that is not an issue).

7.3. Reference Collection

The tool is able to read plain YAML-files in the workspace, read generated output by Kustomize, and fetch objects from a running cluster. The tools does not parse the output of *Helm Charts*. Further research and development is needed to conclude if that is such a feature is feasible.

- **Cluster Scanning:** KRH use the user's current context found in the kube_config to call the Kubernetes Cluster and collects the names of the kubernetes objects that way.
- **Workspace Scanning:** KRH travers all files ending with `.yaml` or `.yml` in the current open workspace in VS Code. It collects the `kind`, `name`, and `namespace` of all the Kubernetes objects found in the files.
- **Kustomize Scanning:** KRH travers all files named `kustomization.yml` or `kustomization.yaml` in the current open workspace in VS Code, and builds the kustomize files. All the kubernetes objects generated by `kustomize` is collected. In addition to that the extension will also highlight if the kustomization-file builds or not.

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization  ✓ Kustomize build succeeded
namePrefix: echo-server-
commonLabels:
  app: echo-server
resources:
- deployment.yml
- service.yml

apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization  ✗ Kustomize build failed
namePrefix: echo-server-
commonLabels:
  app: echo-server
resources:
- thisFileDoesNotExist.yml
- service.yml
```

If *Kustomize* is not installed the computer and part of default shells PATH, it will notify the user.

Each time you save a YAML-file in VS Code the extension will update its internal list of Kubernetes objects.

Each of these scanning techniques can be disabled through the CommandLine or through settings.

Kustomize is extra challenging to handle because it is build on the idea of "[bases](#)" and "[overlays](#)". Overlays use a base-configuration and add/delete/modify/override the existing configuration. Multiple overlays can use the same base and a base has no knowledge of the overlays that refer to it. Overlays can be chained to infinity, which means that when the extension builds a kustomization, it doesn't know if this is the final configuration, or one or more overlays override it in a different folder. Traversing the dependency tree of overlays depending on each other is beyond the scope of this Paper. The tool will instead just inform the user of references it finds in all the kustomization layers, and then it is up to the developer to check which layer he/she meant to refer to. The fact that the extension highlights which kustomize-files builds or doesn't build can be very useful in debugging a long chain of kustomize layers.

8. Evaluation

8.1. Results

8.2. Limitations

9. Conclusion

This paper has

open source

Handles templating

Live feedback

danling references

10. References